Testing Momentum Strategies using Python

BACHELOR'S THESIS

SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENTS FOR THE DEGREE OF
BACHELOR OF ARTS IN BANKING AND FINANCE

AUTHOR
SASKIA SENN
RAINACKERSTRASSE 6, 8953 DIETIKON
19-734-664
SASKIA.SENN2@UZH.CH

CHAIR OF
PROF. DR. MARKUS LEIPPOLD
PROFESSOR OF FINANCIAL ENGINEERING
DEPARTMENT OF BANKING AND FINANCE
UNIVERSITY OF ZURICH

SUPERVISOR
DR. BENJAMIN WILDING
DEPARTMENT OF BANKING AND FINANCE
UNIVERSITY OF ZURICH

DATE OF SUBMISSION: MARCH 9, 2023

Abstract

The aim of this thesis is to develop an automated correction tool using the Python programming language to efficiently correct the *Involving Activity 3* in the course *Asset Management: Investments*. The exercise requires students to create two momentum strategies based on historical stock prices of 18 stocks using varying look-back and holding periods. The tool is designed to be highly flexible in terms of input data, loock-back, and holding periods, enabling the momentum strategies to be effectively tested and compared to a buy-and-hold strategy. The tool offers a powerful approach for correcting the *Involving Activity 3* leading to faster processing times and minimized errors compared to manual correction methods.

I

Institut für Banking und Finance

Universität Zürich
Institut für Banking und Finance
Plattenstrasse 14
CH-8032 Zürich
Telefon +41 44 634 29 51
www.bf.uzh.ch

Bachelorarbeitssauftrag für:

Saskia Senn
Oberwiesenweg 26
5436 Würenlos

Prof. Dr. Markus Leippold

Prof. Dr. Markus Leippold
Professor of Financial Engineering

Zürich, 21. Juli 2022

Auftragserteilung für die Bachelorarbeit

Sehr geehrte Frau Senn

Mit Freude habe ich erfahren, dass Sie bei mir die Bachelorarbeit schreiben möchten. Hiermit erteile ich Ihnen dafür den folgenden Auftrag:

Titel: Testing Momentum Strategies using Python

Zielsetzung:
In the first, short part of your thesis, develop the theoretical foundations as well as an overview of the existing scientific literature - especially empirical studies - on momentum strategies.

In the second part, develop a tool in Python that can be used to test momentum strategies. Make sure that the tool is designed to be flexible with respect to input data and lockback and holding periods. Comment your code in detail.

Finish your work with a conclusion.

Ablauf Ihrer Arbeit

- Sie erhalten die Möglichkeit, eine allfällig erstellte Disposition mit Ihrer Betreuungsperson am Institut zu besprechen. Machen Sie dazu mit Ihrer Betreuungsperson einen Termin aus und senden Sie dazu Ihre Disposition.

Bitte beachten Sie folgende formale Kriterien:

- Sie verfassen die Arbeit in deutscher Sprache.
- Qualität geht vor Quantität. Achten Sie auf kompakte Schreibweise; verzichten Sie auf lange und unnötige Ausführungen und kommen Sie rasch zum Wesentlichen. Ihre Arbeit (exkl. Verzeichnissen und Anhang) sollte maximal 40 Seiten umfassen.
- Achten Sie bitte auf korrekte, fehlerfreie Sprache und einen wissenschaftlichen, knappen, aber flüssigen Schreibstil. Achten Sie auch auf korrekte Zitierweise.
- Verwenden Sie ausreichend beschriebene Graphiken. Graphiken und Tabellen sollten selbsterklärend und unabhängig von der Arbeit verständlich sein.
- Bauen Sie Ihre Arbeit wie folgt auf:
    - Titelblatt
    - Abstract nach dem Vorbild wissenschaftlicher Zeitschriften (max. 100 Worte)
    - Vorliegende Auftragserteilung
    - Inhaltsverzeichnis
    - Hauptteil
- Lade auf OLAT ein ZIP-File mit folgenden Dateien:
    - Titelblatt der Arbeit (PDF-Dokument)
    - Executive Summary (PDF-Dokument): Zusammenfassung auf max. 3 Seiten
    - Gesamte Arbeit (LaTex- oder Word-File sowie PDF-Dokument)
    - Gesamtes verwendetes Datenmaterial
    - Computer-Codes zur Replikation Ihrer Ergebnisse
    - Elektronisch gespeicherte Referenzen (Papers in PDF-Format).

Bei allfälligen Fragen wenden Sie sich an Benjamin Wilding, Institut für Banking und Finance, Email: benjamin.wilding@bf.uzh.ch

Gute Arbeit und viel Erfolg!

Freundliche Grüsse

Markus Leippold
Professor of Financial Engineering
Universität Zürich
Institut für Banking und Finance

Contents

List of Tables

List of Figures

1 Introduction

"The trend is your friend", is a famous quote when it comes to momentum investing. Momentum investing is a strategy that involves buying past winners and selling past losers (Jegadeesh and Titman (1993)). The idea of momentum investing dates back to Levy (1967) who initially considered stocks that historically (called look-back period) have been relatively strong tend to remain relatively strong for a significant period of time (called holding period). This idea challenges one of the key concepts in modern finance, which is the Efficient Market Hypothesis. The Efficient Market Hypothesis states that capital markets are efficient and that the price of an asset already reflects all available information (Fama (1970)). Thus, implying that it shall not be possible to gain superior returns over the market.

According to Fama (1970), a market in which prices always fully reflect all available information is called efficient. There are three stages of an efficient market. Firstly, the weak form focuses on a limited subset of information, namely, past price or return histories (Fama (1970)). Secondly, the semi strong form in which prices seem to efficiently adjust to publicly available information and finally, the strong form which investigates whether certain investors have exclusive access to relevant information for price formation (Fama (1970)). Momentum, therefore, appears to violate the efficient market hypothesis in its weakest form (Ehsani and Linnainmaa (2022)). For the reason that past returns should not predict future returns if asset prices respond to new information immediately and to the right extent (Ehsani and Linnainmaa (2022)).

The term "relative strength" was first introduced by Levy (1967) as an earlier form of momentum. However, momentum gained more attention after the influential work of Jegadeesh and Titman (1993). Their study observed that stocks with higher past returns tend to continue performing well in the future for three to twelve months, while stocks with lower past returns tend to underperform over the next few months. They concluded that investors can earn significant profits by purchasing past winners and selling past losers.

The bachelor course *Asset Management: Investments* at the University of Zurich offers a comprehensive examination of the key components of the investment process, ranging from the identification and assessment of investment instruments to the creation and evaluation of optimal portfolios. The course outlines a strong foundation in asset management principles, focusing on capital market theory, asset allocation, investment strategies, equity, and fixed income instruments, as well as performance measurement. The annual attendance for this lecture exceeds 200 students.

The students[1] of this course voluntarily solve three Excel-based exercises called *Involving Activities* during the semester in order to gain exam points in advance. The individual exercises serve to deepen the material and apply the theory previously stud-

---

[1] The course involves students and executive education participants to solve the relevant exercise but for simplicity reasons those two groups are summarized and only referred to as students from now on.

ied. The *Involving Activity 3* is designed to create a cross-sectional momentum strategy based on 18 given stocks. Thereby, the students build two momentum strategies by choosing their individual look-back and holding period. The first strategy is a long-only strategy, referring to buying stocks with the anticipation of its increase in market value, and the second one a long-short strategy, the latter further implies selling stocks with the anticipation of a decrease in market value. At the end of the *Involving Activity 3* they will compare their results with a buy-and-hold strategy.

Until now, the *Involving Activity 3* has lacked a clear structure, enabling a very individual interpretation by every student resulting in various solutions. This had the effect of every exercise having to be corrected manually resulting in a huge effort and time exposure for the student assistants. Thus, the following thesis aspires to update the *Involving Activity 3* among two different aspects. The exercise should, on one hand, remain to be individual for every student and, on the other hand, allow for an automated correction by structuring and standardizing the exercise more. Within the scope of this thesis, a correction tool has been developed and written in the programming language Python to evaluate the momentum strategies exercise of the students. The tool is designed to allow for flexibility in the input data, look-back and holding periods. The more structured exercise enhances the students' understanding and promotes successful outcomes. Additionally, the integration of automated correction represents an efficient solution for grading, leading to significant time savings and greater objectivity in evaluation. The tool also promotes comparability among students, contributing to more informed decision-making.

The thesis is divided into three sections. The following section outlines an overview of influential literature on the cross-sectional and time-series momentum. Section 3 describes the *Involving Activity 3* and the Python code in detail. Lastly, Section 4 outlines the conclusion and suggests an outlook for further implementations.

2   Literature Review

As mentioned before, the students were tasked in the *Involving Activity 3* to create their own momentum strategy based on historical stock prices. To achieve this, it is crucial to have a solid understanding of the key research studies that have investigated the momentum strategy's effectiveness. Therefore, this section provides a comprehensive overview of the primary influential studies related to the momentum strategy. The momentum strategy is supported by various empirical studies and there exist many different behavioral explanations for why momentum investing leads to abnormal returns.

Moskowitz et al. (2012) showed that positive momentum that partially reverses in the long-term could be in line with the theory of initial under-reaction and delayed over-reaction. According to Grinblatt and Han (2002), under-reaction can be caused by the disposition effect, selling winners too early and riding losers too long which may be driving momentum in stock returns. Hong and Stein (1999) showed that if information spreads slowly throughout the population, prices may initially under-react leading to potential profits for momentum traders who engage in trend-chasing strategies. Over-reaction can be posed, according to Hoitash and Krishnan (2007), by herding behavior. Another explanation by Barberis et al. (1998) is the representativeness heuristic that investors might conclude that the past history of an underlying is representative of its potential growth.

There are two types of momentum strategy: cross-sectional and time-series momentum. Most literature focuses on the relative performance of securities, on a cross-section basis, thereby finding that securities that recently outperformed their peers over the past three to twelve months continue to outperform their peers on average over the next three to twelve months (Jegadeesh and Titman (1993)). Accordingly, the theoretical foundations of the cross-sectional momentum are outlined first. Afterwards, literature on the time-series momentum is summarized. Rather than focusing on the relative returns of securities in the cross-section, time-series momentum focuses exclusively on a security's own past return (Moskowitz et al. (2012)).

2.1   Cross-sectional Momentum

Jegadeesh and Titman (1993) found that trading strategies that involve buying past winners and selling past losers realize significant abnormal returns between 1965 to 1989. For instance, a strategy that selected stocks based on their past six-month returns and held them for six months resulted in an average compounded excess return of 12.01% per year. Lewellen (2002) found that the momentum effect is significant for individual stocks and portfolios for seven to nine months after formation but quickly diminishes and turns to contrarian profits. This is in line with the findings of Jegadeesh and Titman (1993) and Moskowitz and Grinblatt (1999). However, Bhattacharya et al. (2015) used the identical dataset over the 1999 to 2012 timeframe, and reported that momentum strategies, as outlined by Jegadeesh and Titman (1993), failed to generate abnormal

returns.

Rouwenhorst (1998) analyzed return continuation in twelve European countries between 1980 and 1995 and found that a diversified portfolio of past winners outperformed a portfolio of past losers by approximately 1% per month. Rouwenhorst (1998) also validated the presence of momentum strategies in emerging markets.

Chan et al. (2000) investigated the profitability of momentum strategies for individual international equity stock indices by analyzing past returns. The findings suggest the existence of significant momentum profits, particularly for short holding periods of less than four weeks. Griffin et al. (2005) expanded the research on 40 markets and found evidence for high profits in momentum strategy. They showed that momentum strategies generate substantial profits in a variety of markets. According to Kang et al. (2002), there were statistically significant abnormal profits associated with momentum strategies across eight distinct look-back and holding periods on the China stock market from 1993 to 2000. It was observed by Chan et al. (2019) that strategies focusing on past returns lead to significant profits over a six to twelve month timeframe, based on all the stocks listed on the NYSE, Amex, and Nasdaq. For instance, sorting stocks according to their prior six-month return generated a return spread of 8.8 percentage points during the subsequent six months. Chui et al. (2010) suggested that momentum strategies are found to be profitable in most regions around the world, both in developed and emerging markets, during the period 1980 to 2005. They also noted that the profitability of momentum strategies varies across regions, with some regions exhibiting higher levels of profitability than others.

Naughton et al. (2008) examined momentum trading strategies for Shanghai Stock Exchange equities and found significant profits between 1995 to 2005, taking into account trading volume and utilizing different formation, and holding periods. Contrary, Hameed and Kusnadi (2002) reported limited evidence to support the presence of momentum in six emerging Asian stock markets over holding periods of three to twelve months. Hameed and Kusnadi (2002) found that these trading strategies consistently generate insignificant profits, indicating a lack of momentum effect in the Asian markets studied.

Jostova et al. (2019) documented strong evidence of momentum profitability in U.S. corporate bonds. Based on their dataset from 1991 to 2011, they revealed that past six-month winners outperform losers by an average of 59 basis points per month if held for six months. Additionally, momentum strategies were found to be profitable across look-back and holding periods of three, nine, and twelve months. Beyhaghi and Ehsani (2016) also documented momentum profitability in loans of U.S. firms. Grinblatt et al. (1995) found that mutual fund momentum strategies were able to generate significant excess performance.

A concurrent study by Carhart (1997) discovered that purchasing mutual funds that were top-performers in the previous year while selling those that performed poorly can potentially result in higher returns. Miffre and Rallis (2007) found that momentum strategies applied to 31 US commodity futures contracts resulted in profitable returns

over horizons that vary between one to twelve months in commodity futures markets.

Menkhoff et al. (2012) found significant excess returns of momentum strategies between past winning and losing currencies in the foreign exchange market. According to Bhojraj and Swaminathan (2006), analyzing 38 country indices, winners outperform losers in the first three to twelve months after portfolio formation, but underperform losers in the following two years. Moreover, Okunev and White (2003) showed examining foreign exchange markets for Australia, Canada, France, Germany, Japan, Switzerland, the U.K., and the U.S. from January 1975 through June 2000 that the use of momentum strategy generated excess returns.

Asness et al. (2013) conducted a comprehensive analyzis of eight distinct markets and asset classes which included individual stocks in the United States, the United Kingdom, continental Europe, and Japan, as well as country equity index futures, government bonds, currencies, and commodity futures. They revealed compelling evidence of momentum return premia across all the markets studied, indicating a pervasive and robust presence of momentum in the global financial markets.

2.2 Time-series Momentum

Time-series momentum, as noted by Moskowitz et al. (2012), is related to but still different from the cross-sectional momentum phenomenon in finance literature. The cross-sectional momentum approach involves selecting past winners and avoiding past losers, while time-series momentum focuses on trend-following based on a stock's past price movement. Predominantly, studies have mainly focused on futures markets and managed futures funds with regard to time-series momentum (Moskowitz et al. (2012), Hurst et al. (2017), Baltas and Kosowski (2013), Ham et al. (2019), Kim et al. (2016), Jin et al. (2019)).

As one of the first, Moskowitz et al. (2012) focused purely on a security's own past return and documented significant time-series momentum across very different asset classes and markets. Moreover, they found that the past twelve-month excess return of each instrument is a positive predictor of its future return. These findings are valid across a number of subsamples, look-back periods, and holding periods. Also, Moskowitz et al. (2012) showed that the existence and significance of time-series momentum is robust across horizons and asset classes, particularly when the look-back and holding periods are twelve months or less. In addition to Moskowitz et al. (2012), Baltas and Kosowski (2013) provided the first concrete empirical evidence of time series momentum using a broad daily data set of futures contracts. They showed that the time-series momentum strategy generates a statistically and economically significant mean return and alpha. Ham et al. (2019) documented significant time-series momentum profits in ten commodity futures from 2006 to 2018 in the Chinese market. Moreover, Hurst et al. (2017) reported that time-series momentum has been consistently profitable throughout the past 137 years. Specifically, they constructed combinations of one-month, three-month

and twelve-month time-series momentum strategies from January 1880 to December 2016. Furthermore, they found that a trend-following strategy performed relatively similarly across a variety of economic environments and provided significant diversification benefits to a traditional allocation.

Kim et al. (2016) re-examined the time-series momentum strategy using 55 liquid futures contracts. They found that time-series momentum only significantly outperforms a buy-and-hold strategy during the 1985 to 2009 period if the time-series momentum employed volatility scaling. Also, Huang et al. (2020), using the same data set as Moskowitz et al. (2012), found that the evidence on time-series momentum is weak and that time-series momentum across the global asset classes appears questionable.

Comparatively, less attention has been devoted to the most conventional of asset classes: common stocks. Lim et al. (2018) document strong time-series momentum effects in individual stocks in the US markets from 1927 to 2017. Moreover, time-series momentum is not specific to sub-periods, firm sizes, formation- and holding-period lengths, or geographic markets. Also, Georgopoulou and Wang (2017) documented a significant time-series momentum effect that is consistent and robust across global equity and commodity markets from 1969 to 2015. The findings of Shi and Zhou (2017) indicated that there is a time-series momentum effect in the short run and a contrarian effect in the long run in the Chinese stock market. The performances of the time series momentum and contrarian strategies are highly dependent on the look-back and holding periods and firm-specific characteristics. Recently, Fang et al. (2021) show that profits may not be generally available for time-series momentum, considering 2.25 million monthly returns on over 20,000 US individual stocks from 1986 to 2017.

Later, other studies document inter-day time-series momentum effects. Gao et al. (2018) showed that the market return in the first half hour of the trading day predicts the market return in the last half hour. They showed an intra-day momentum pattern based on high frequency S&P 500 exchange-traded fund (ETF) data from 1993 to 2013. Jin et al. (2019) reported a similar intra-day time-series momentum across four Chinese Commodity future contracts from 2002 to 2013.

He and Li (2015) showed that the performance of momentum strategy is determined by both time horizon and the market dominance of momentum traders. Specifically, when the momentum traders dominate the market, the momentum strategy is profitable for a short time horizon and unprofitable when the time horizon is long.

D'Souza et al. (2016) stated the significant profitability of time-series momentum strategies in individual stocks in the US markets from 1927 to 2014 and in international markets since 1975. Moreover, the profitability of the strategy was robust for 16 different combinations of formation and holding periods, different benchmarks, and weighting systems in up and down markets. The outcomes of Zakamulin and Giner (2019) suggest that the long-only time-series momentum strategy is profitable, while the long-short one is not. In their study they included the monthly total returns on the Standard and Poor's Composite stock price index, and the sample period begins in January 1857 and

ends in December 2018.

Moreover, Cheemaa et al. (2018) found evidence that time-series momentum returns exceed cross-sectional momentum returns. Due to its active position, which involves taking a net long or short position based on the direction of the market, while cross-sectional strategy is a zero-cost strategy that does not depend on the market state. Goyal and Jegadeesh (2017) adjusted for the net long positions of the time-series momentum strategy and found that, on average, the excess returns of cross-sectional and time-series momentum strategies are generally equal.

3 Correction Tool

In the first section, the task and the Excel file serving as a template for the students to solve is described. Also, the input data utilized is pictured. Then a detailed explanation of the Python code, used to correct the *Involving Activity 3*, is given.

3.1 Task and Excel file

The purpose of the Excel-based *Involving Activity 3* is to compare various risk and return measures of the student's long-only and long-short momentum strategies against a buy-and-hold strategy. Thereby, the students were asked to create a cross-sectional momentum strategy based on their chosen look-back and holding period. The *Involving Activity 3* consists of one Excel file featuring a total of eight sheets ("Einleitung", "Eingabe der Daten", "Grunddaten", "Berechnung mon. Renditen", "Ranking", "Kauf-& Verkaufsignal", "Monatliche Portfoliorendite" and "Gesamtrendite & SR").

In the first sheet "Einleitung", the students had to state their name and matriculation number and are provided with important information regarding the *Involving Activity 3*. In the sheet "Eingabe der Daten" the students did choose their look-back and holding period which could range from three to twelve months as per the students' preference. These time periods are commonly studied and utilized in the literature. Additionally, the students could choose if they rather wanted to base the ranking on the arithmetic or geometric mean. The students are also expected to determine the number of shares to be bought or sold per month for both their long-only and long-short portfolios. The given data is provided in the sheet "Grunddaten" (see Section 3.2).

Following, as a first step, in the sheet "Berechnung mon. Renditen", the students calculated the monthly returns for the entire time period, taking into account all the stocks. In the sheet "Ranking", the students first calculated the average monthly return of 18 stocks based on their chosen look-back period. After calculating the returns, the ranking could be determined based on the total return. From the established ranking, the students invested equally weighted (1/N) in the number of stocks according to their long-only and long-short strategy, respectively. As a result, in the sheet "Kauf- & Verkaufsignal" they calculated the total return over their holding period. In the sheet "Monatliche Portfoliorenditen" the students calculated the monthly portfolio return for the three strategies. As a last step, the students determined in the sheet "Gesamtrendite & SR" the total return, different risk masses, and the sharpe ratio of the three strategies.

3.2 Data

The students were provided with monthly total return values of 18 stocks covering the time period from January 2002 to November 2022 (see Table 1). Thus, the students would invest equally weighted in the number of stocks according to their long-only or long-short strategy, respectively, and calculate the total return over their holding period.

Table 1: Given Data: Extract of the monthly total return values that have been provided to the students

Extract of the sheet "Grunddaten" of monthly total return values of the 18 stocks that have been provided to the students. Note: for illustration purposes only five out of 18 stocks and four data points out of 20 years are presented.

Datum	ABB	CS	GEBERIT	GIVAUDAN	HOLCIM
01.01.2002	2762.73	1565.94	118.71	101.15	1670.80
01.02.2002	2978.23	1726.93	116.67	100.55	1730.38
01.03.2002	2762.73	1580.58	113.05	109.29	1757.79
01.04.2002	2336.04	1414.72	113.05	111.28	1770.90

3.3 Implementation

The Python code is designed to perform the thorough correction process of the *Involving Activity 3* for all students[2], substituting the manual correction in order to speed up the correction process. The aforementioned individual choice of some parameters by the student posed a basic requirement to the correction tool. Therefore, the Python code must be able to adapt flexibly to these input parameters in order to correct the respective *Involving Activity 3* of the student. Before diving into the detailed explanation of the code using flowcharts, a brief overview of the code structure will be provided by referring to the flowcharts to help understand the code better.

3.3.1 Overview

The flowchart in Figure 1 is a simple outline of the Python code structure. The code structure is divided into two Python files: the directory and the correction file.

The first step involves the setup of the directory, in the directory file, to access the submitted student's files. By setting the path to the folder where all the submitted files are stored, the code loops through all these folders in the specified path and searches for Excel files within a subdirectory named "2_submissions". If an Excel file is found, it then calls the `correction` function from the second Python file to correct the file. The setup of the directory is further described in Section 3.3.2.

Step two consists of the correction Python file, outlined in Section 3.3.3, that applies the correction process for each student and the *IA Output*.

In the correction file, as the third step, two main functions are defined (see Section 3.3.4). First, the function `load_workbook_range` is defined (see Figure 4). Every worksheet of the *Involving Activity 3*, as described above, will be loaded into a dataframe in Python and has the same format. Therefore, using this function, data with a specified range in a worksheet can be loaded and a dataframe is returned. It

---

[2] The only difference to the executive education participants is that they do not have a matriculation number, and the path to the submitted files has to be changed. For simplicity reasons, only the correction tool for the students is described and analyzed.

can optionally include headers and an index column with a specified name. After that, the function `correction(filenames)`, as shown in Figure 5, is defined. The `correction(filenames)` consists of several additional steps.

The `correction(filenames)` function is called and as the fourth step, the student file and the empty solution file are loaded into dataframes (see Section 3.3.5). Figure 6 illustrates how the student file is loaded into dataframes. This step also involves saving important input parameters and defining variables. Table 2 shows the data which is used to calculate the solution of the monthly return. All other calculations are based on the student's value in order to take consequential errors into account.

As the fifth step, the student file will be compared to the solution file (see Section 3.3.6). For that, the solution is calculated and the difference in the two dataframes is built for assigning points for each exercise.

In the sixth and last step, outlined in Section 3.3.7, the *IA Output* is generated which contains information about the student, the points they achieved and if they passed or failed the exercise. This information can then be extruded to perform the mass evaluation on OLAT. Also, the Flowchart 12 illustrates that process in detail.
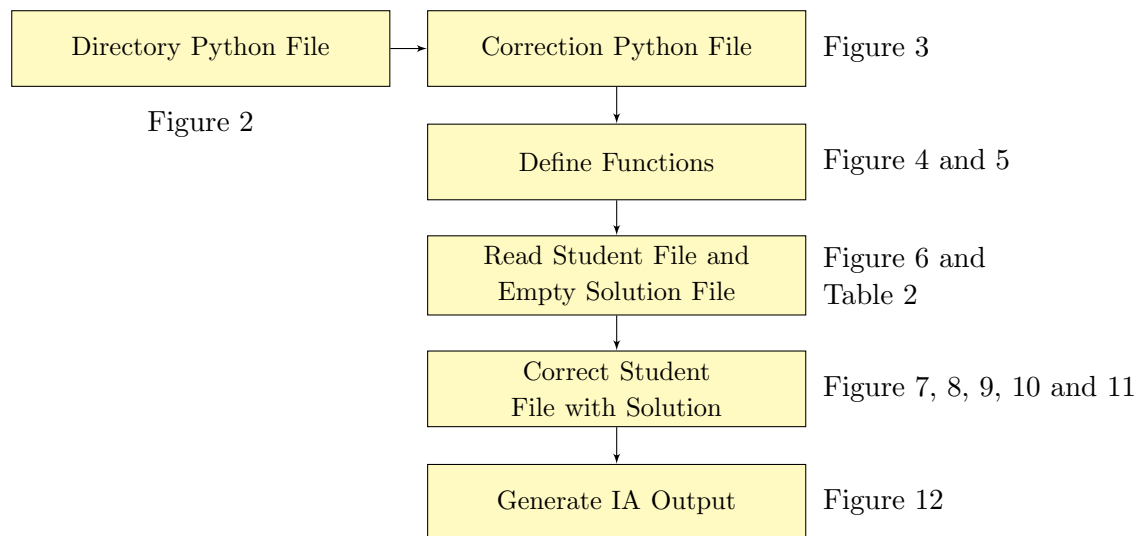
Directory Python File	Correction Python File	Figure 3
Figure 2	Define Functions	Figure 4 and 5
	Read Student File and Empty Solution File	Figure 6 and Table 2
	Correct Student File with Solution	Figure 7, 8, 9, 10 and 11
	Generate IA Output	Figure 12

Figure 1: Flowchart Overview
Overview of the Python code with refering to the relevant sections and figures that describe the process in detail.

Flowcharts use various shapes to represent different elements. Each flowchart starts with a red rounded rectangle symbol representing the start of the process. Purple parallelograms are used to represent input or output, while orange rectangles describe the process being performed. Green diamonds indicate a decision point in the flowchart, while orange circles are connectors for `for-loops`. Lines are used to connecting the various shapes and show their relationships and dependencies. Finally, a red rounded rectangle symbol is used to signify the end of the flowchart.

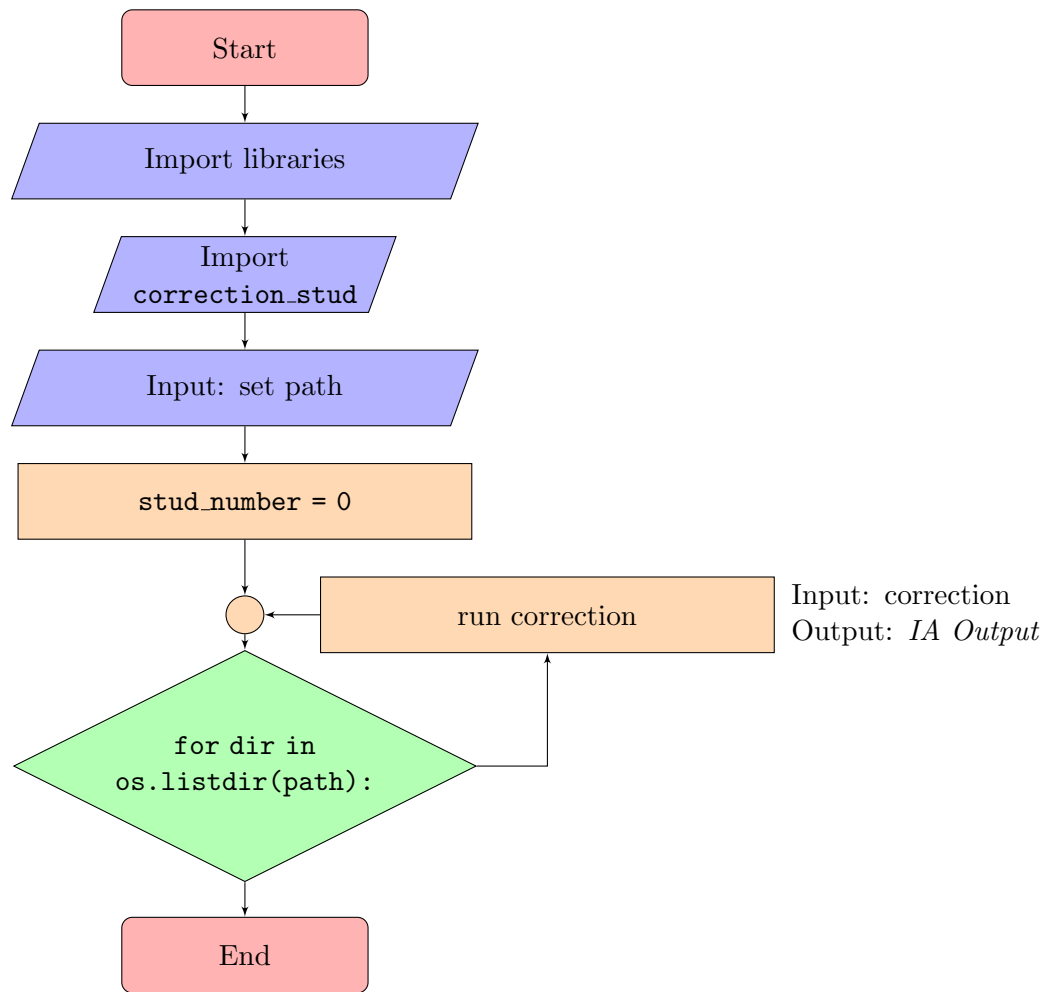### 3.3.2 Directory Python File

Figure 2: Flowchart dir_stud.py

The main part of the Directory Flowchart (see Figure 2) is the `for loop` iterating through each student file.

To enable this `for loop`, the `correction_stud` (see Figure 3) needs to be imported which contains the second Python file with the correction and *IA Output* in it. Then the file path, where the student's submitted files are stored, is defined and initializes a counter for the number of students processed, `stud_number`, to zero.

Next, the code uses a `for loop` that iterates through each directory in the specified path using the `os.listdir` function. The `os.listdir(path)` function retrieves a list of all the directories in the specified path. The variable `dir` will hold the name of each directory in the list in each iteration of the `loop`. With this `for loop`, access is granted to the submission folder which contains the submitted Excel sheets of the students. In the code, this is represented with the filenames. Therefore, in each iteration of the loop, the code accesses the *Involving Activity 3* of one of the student's directories within the file path. This allows the code to perform the correction process for each student's

Involving Activity 3 individually, instead of processing all the submissions in a single run. For each directory, the code uses the `glob` library to search for Excel files with a ".xlsx" extension. The `os.path.join` function is used to construct the full path to each student's submissions directory. The `if` statement checks if the filenames list is empty or not. If filenames is empty, meaning that no Excel files were found in the current directory, then the `correction_stud.correction` function is not called, and the code continues to the next directory in the `os.listdir` iteration.

The code keeps track of the number of students processed by increasing the `stud_number` variable each time it processes a directory.

The code then calls the `correction_stud.correction` function with the first Excel file in the filenames list as an argument. This function is defined in the `correction_stud.py` and performs the correction on the *Involving Activities*. This function is described in detail in the following paragraph (see Figure 3).
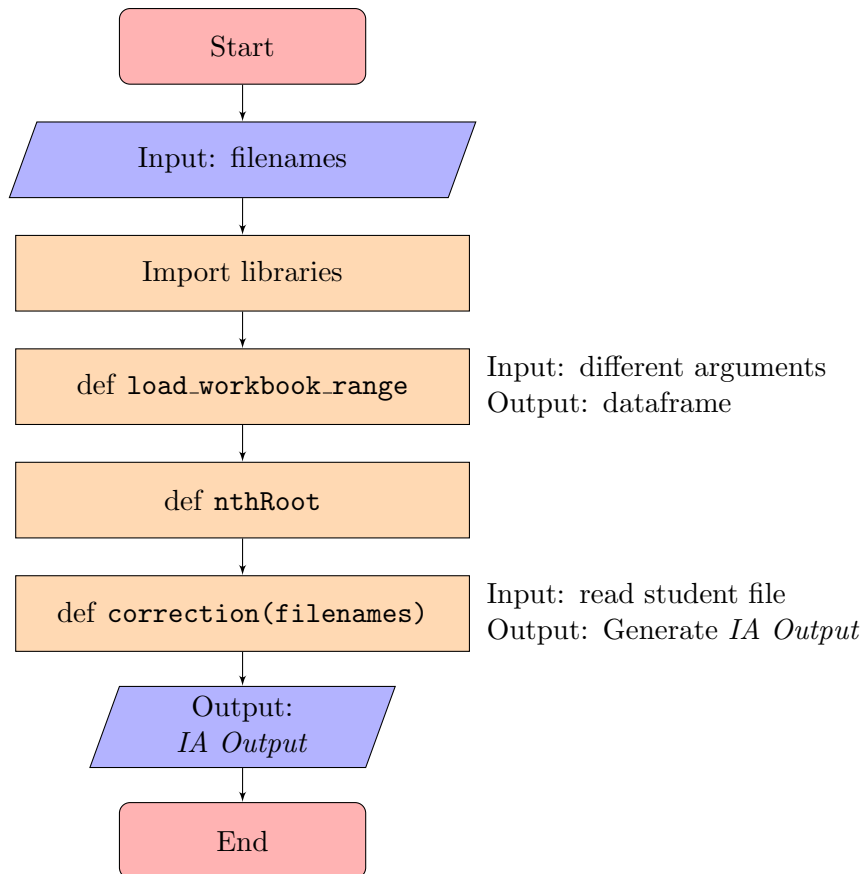
3.3.3 Correction Python File

Figure 3: Flowchart corr_stud.py

The correction process within the code involves the use of dataframes for efficient processing of data. Given the 18 shares and a 20-year time period, dataframes provide a convenient and scalable method for organizing and analyzing the data in Python.

12

First, the code defines the function `load_workbook_range` that can be appreciated in Figure 4. This function is used to load specific data ranges from an Excel workbook.

Additionally, the function `nthRoot` is used to find the "n^{th}" root of a given value "x".

The core correction process is performed by the function `correction(filenames)` which is described in detail in Figure 5.

3.3.4 Define Functions

The `load_workbook_range(range_string, ws, with_header=True, with_index=False, index_name=None)` function is a function that loads data from an Excel sheet into a dataframe.

It takes the following parameters:

- `range_string`: a string specifying the range of cells in the worksheet to be loaded into the dataframe. The format of the string should be in the form of "A1:Z100"

- `ws`: a worksheet object representing the worksheet in the workbook.

- `with_header`: a boolean value indicating whether the first row of the data should be used as the header for the dataframe.

- `with_index`: a boolean value indicating whether the data should be set as the index of the dataframe.

- `index_name`: a string specifying the name of the index if `with_index` is set to True. If `with_index` is False, this parameter is ignored.
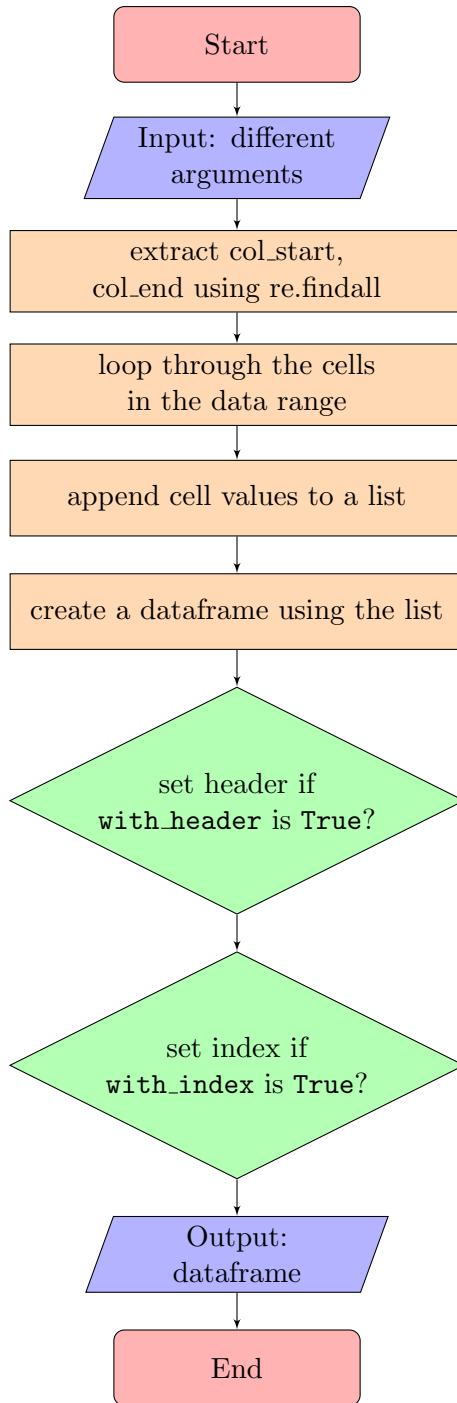
Figure 4: Flowchart `def load_workbook_range()`

The function first uses the `re` module to extract the start and end columns of the data range. It then loops through the cells in the data range and appends the cell values as a list to `data_rows`.

Next, a dataframe is created from the `data_rows` and the columns are set to the column names obtained from the `range_string` using the `get_column_interval` function.

If `with_header` is `True`, the first row of the dataframe is set as the header. If the `with_index` flag is `True` and the `index_name` is not `None`, the dataframe is given an

index using the `index_name` as the label.

The `print` statement then prints the columns of the dataframe. Finally, the dataframe is returned.
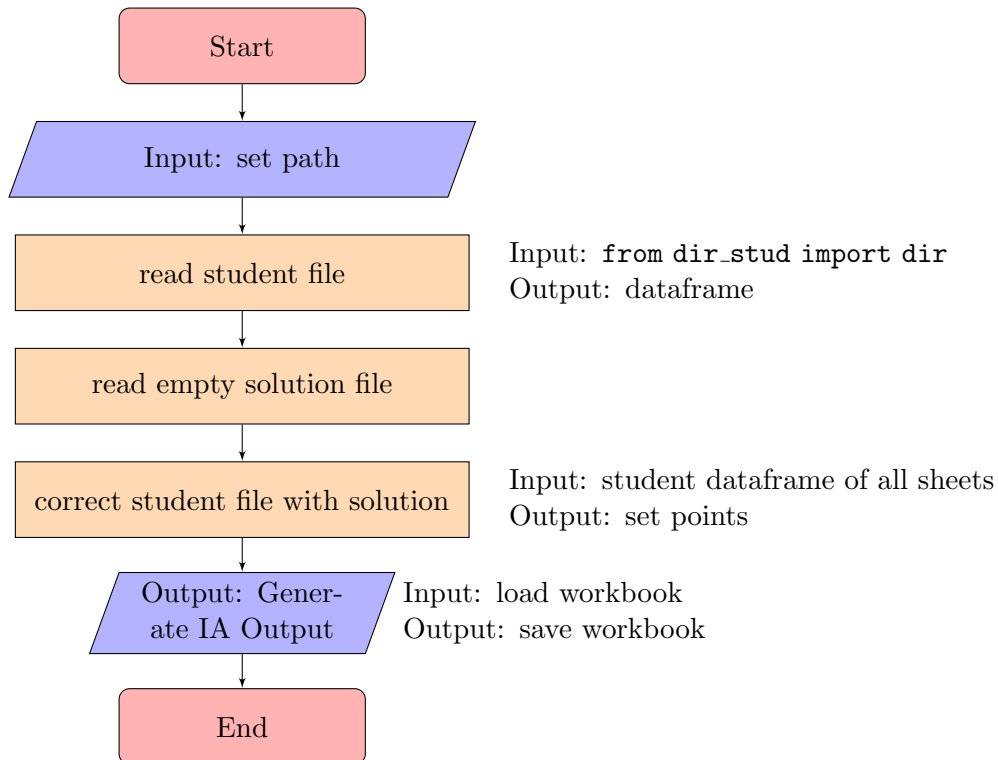

Figure 5: Flowchart `def correction(filenames)`

The function `correction(filenames)` as shown in Figure 5 outlines the correction process for each student's *Involving Activity 3*.

First, the path where the solution file and the empty *IA Output* files are stored, needs to be adjusted. Then the process consists of several steps each of them described in further detail in the following figures.

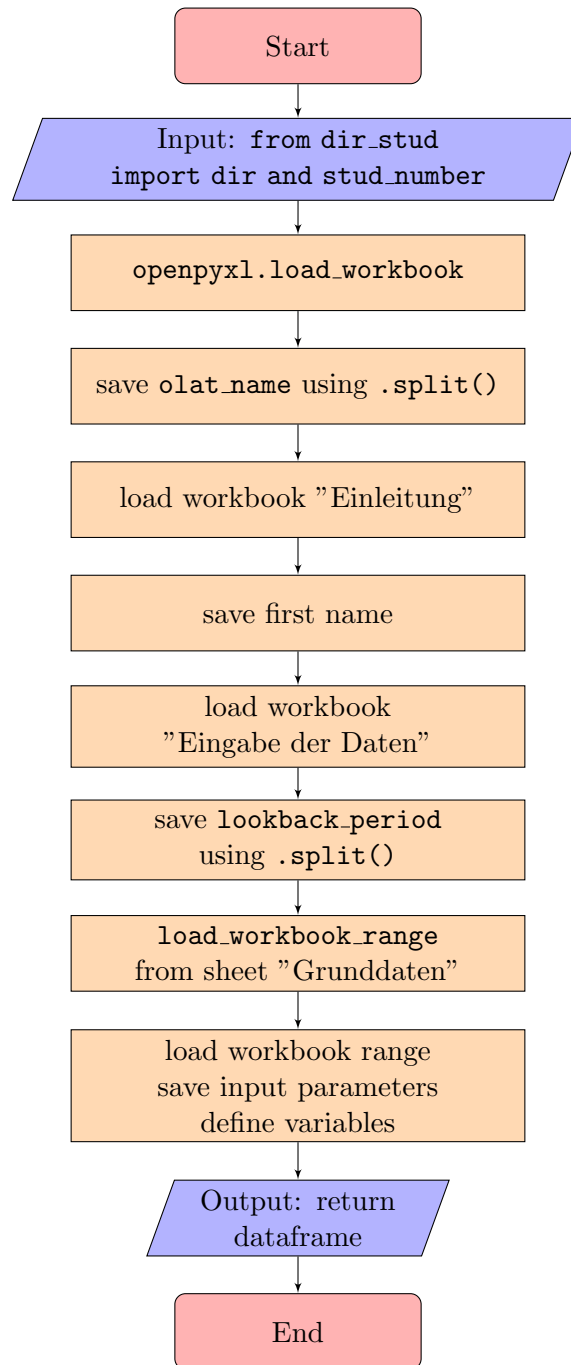### 3.3.5 Read Student File and Empty Solution File

Figure 6: Flowchart Read Student file

The first step involves reading the student's file illustrated in Figure 6. For that the code, as shown in Figure 6, firstly imports two variables, `dir` and `stud_number` from the `dir_stud` module. The "from" keyword is used to specify the module from which the variables are being imported. After this line has been executed, the `stud_number` and `dir` variables are accessible in the current scope and can be used just like any other variables.

Next, the code uses the `openpyxl` library to load an Excel file into memory and create a workbook object. The `load_workbook` function is called with the filenames argument and the `data_only` argument is set to `True`, which means that any formulas in the workbook will be replaced with their calculated values. The resulting workbook is stored in the `wb_stud` variable which can be used to access and manipulate the data in the workbook.

The code then splits the value stored in the dir variable into separate parts using the `split` method which separates a string into a list of substrings based on a specified delimiter. The delimiter defined is: "_". The resulting list of substrings is then indexed using square bracket notation with `-1` as the index. The `-1` index retrieves the last element in the list. This value, after being split by "_", is assigned to the `olat_name` variable.

After loading the workbook, the code accesses the specific worksheet in the workbook, called "Einleitung". This worksheet is loaded into the `ws_einleitung` variable. Then, the code searches for a specific cell in the worksheet, at row 11 and column 5, using the `.cell()` method. The value of this cell is retrieved using the value attribute and stored in the `first_name` variable. It is important to note that in the `.cell()` method, the first row and first column in the worksheet are both numbered 1, not 0. Thus, row 11 and column 5 refer to the 11th row and 5th column in the worksheet, respectively. The process will be repeated for other relevant variables. After that, the maximal reachable points in each exercise are saved as a variable.

Then, the next worksheet "Eingabe der Daten" is loaded into the `ws_eingabe_der_daten`. The code again accesses a specific cell in this worksheet where the students could choose their look-back period. The code takes the value stored in the `lookback_period_month` variable and splits it into separate parts using the `split` method. The delimiter used here is a space character: " ". The resulting list of substrings is then indexed using square bracket notation with 0 as the index. The 0 index retrieves the first element in the list. The value of the first element in the list, after being split by a space character, is then converted to an integer using the `int` function and assigned to the `lookback_period` variable. For example, if the value of `lookback_period_month` is "12 months", then after splitting the value using split(" ", 1), the list would be ["12", "months"]. The value of "12" would be converted to an integer using int("12"), which would result in the integer value 12. This integer value would then be assigned to the `lookback_period` variable. The same process goes with the holding period. Some more variables are defined.

The line, `ws_grunddaten_stud = wb_stud["Grunddaten"]` accesses the worksheet named `"Grunddaten"` in the workbook object stored in the `wb_stud` variable. The result of this expression is a worksheet object, which is assigned to the `ws_grunddaten_stud` variable. The line, `df_grunddaten_stud = load_workbook_range("C10:U261", ws_grunddaten_stud, with_index=True, index_name="Datum ")` uses the `load_workbook_range` function (see Figure 4) to load data from a specified range of cells in the worksheet into a dataframe. The first argument to the `load_workbook_range function, "C10:U261",`

17

is a string specifying the range of cells to load using the A1 notation. This string specifies that the data should be loaded from the cells in columns C to U and rows 10 to 261. The second argument to the `load_workbook_range` function, `ws_grunddaten_stud`, is the worksheet object from which the data should be taken. The `with_index` argument is set to `True` implying that the first row of the specified range of cells will be used as the index (row labels) of the resulting dataframe. The `index_name` argument is set to `"Datum "` which specifies the name to use for the index (row labels) of the resulting dataframe. The result of the `load_workbook_range` function call is assigned to the `df_grunddaten_stud` variable, which will now contain a dataframe with the loaded data. The remaining worksheets will be loaded into a dataframe the same way.

After loading the studentfile into dataframes in Python the empty solution file with the same workbook range is loaded into Python. The time period in the empty solution file in the sheet "Grunddaten" had to be adjusted by starting the time period already at first December 2001 instead of first January 2002 (see Table 2). Because of this adjustment the `.pct()` function could be used to calculate the monthly return starting in January 2002. Otherwise the first monthly return would be stated in February 2002. This was done to ensure that the use of the percentage change function would result in a comparable index between the solution values and the student values.

Table 2: Shifted Data: Extract of the monthly total return values for calculating the solution

Extract of the sheet "Grunddaten" of monthly total return values of the 18 stocks that have been used to calculate the solutions. Note: for illustration purposes only five out of 18 stocks and four data points out of 20 years are presented.

Datum	ABB	CS	GEBERIT	GIVAUDAN	HOLCIM
01.12.2001	2762.73	1565.94	118.71	101.15	1670.80
01.01.2002	2978.23	1726.93	116.67	100.55	1730.38
01.02.2002	2762.73	1580.58	113.05	109.29	1757.79
01.03.2002	2336.04	1414.72	113.05	111.28	1770.90

3.3.6 Correct Student File with Solution

Having organized all the necessary data in Python, the file is prepared to start iterating through each worksheet and using it to correct and grade the corresponding student worksheet.
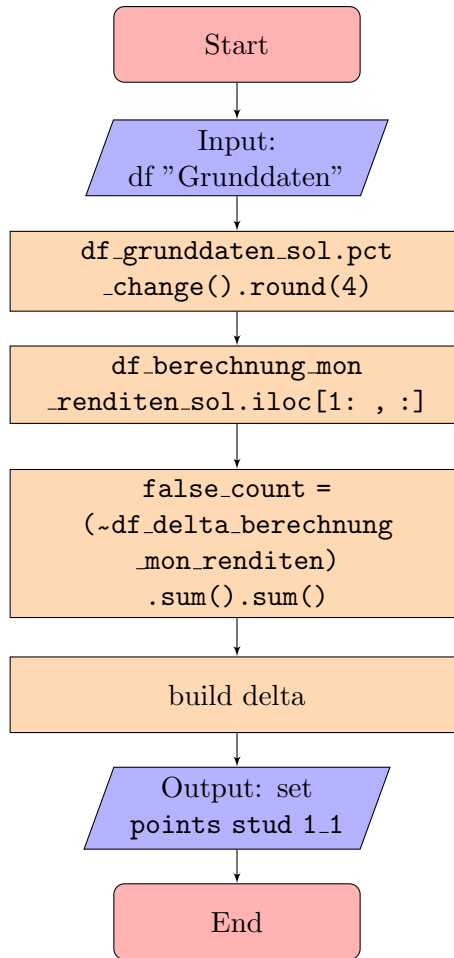
Figure 7: Flowchart Sheet Monatliche Rendite

In the first worksheet, shown in Figure 7, the students had to calculate the monthly return of every share. Firstly, the solution based on the shifted data (see Table 2) is calculated. The `.pct_change().round(4)` function calculates the percentage change of each value in the `df_grunddaten_sol` dataframe and stores the result in a new dataframe `df_berechnung_mon_renditen_sol`. The result is rounded to four decimals. To adjust the solution data, the `.iloc()` function is used to remove the first row from the `df_berechnung_mon_renditen_sol` dataframe. The `.iloc()` method is utilized to access specific rows and columns in a dataframe based on their integer position. The `[1: , :]` syntax specifies that all rows starting from the second row (index 1) and all columns (`:`) should be selected. Thus, this line of code selects all rows in `df_berechnung_mon_renditen_sol` starting from the second row, and all columns.

The function `.astype(float).round(4)` converts the data in the `df_berechnung_mon_renditen_stud` dataframe to the float data type and rounds the result to four decimal places. This step is necessary as the result calculation can differ between Excel and Python.

Next, the difference between the solution and student dataframes is calculated by equating the student dataframe with the solution dataframe. To determine any dispari-

ties, a delta dataframe is created which indicates instances where the difference between the two dataframes was non-zero, thereby expressing a discrepancy between the student's result and the solution. Contrarily, instances with no difference are recorded as `True`.

Then, the amount of False values in the `df_delta_berechnung_mon_renditen` dataframe get counted. The `~` operator inverts the Boolean values in the dataframe so `True` values become `False` and vice versa.

The `sum` method is then used to count the number of `False` values in each column of the dataframe, and the sum method is called again to find the total number of `False` values across all columns. This total is stored in the `false_count` variable. After calculating how many values are false in the student dataframe, the points for this first exercise that the students receive are calculated. The value of `points_stud_1_1` is reduced based on the number of `False` values in the `df_delta_berechnung_mon_renditen` dataframe. Every correct value gives a sub-point. The sub-points add up to the maximum number of points which could be achieved in the exercise. So, for every false value, a sub-point is subtracted and saved the achieved points in the variable `points_stud_1_1`.

The sequence of rounding the values, building the delta, counting the false values, and setting the points for the exercise is the same process in every sheet. Therefore, this process will not again be described in detail in the following extensions.

Then a new dataframe is created by adding one to the monthly return dataframe. This will later be useful for calculating the geometric mean.
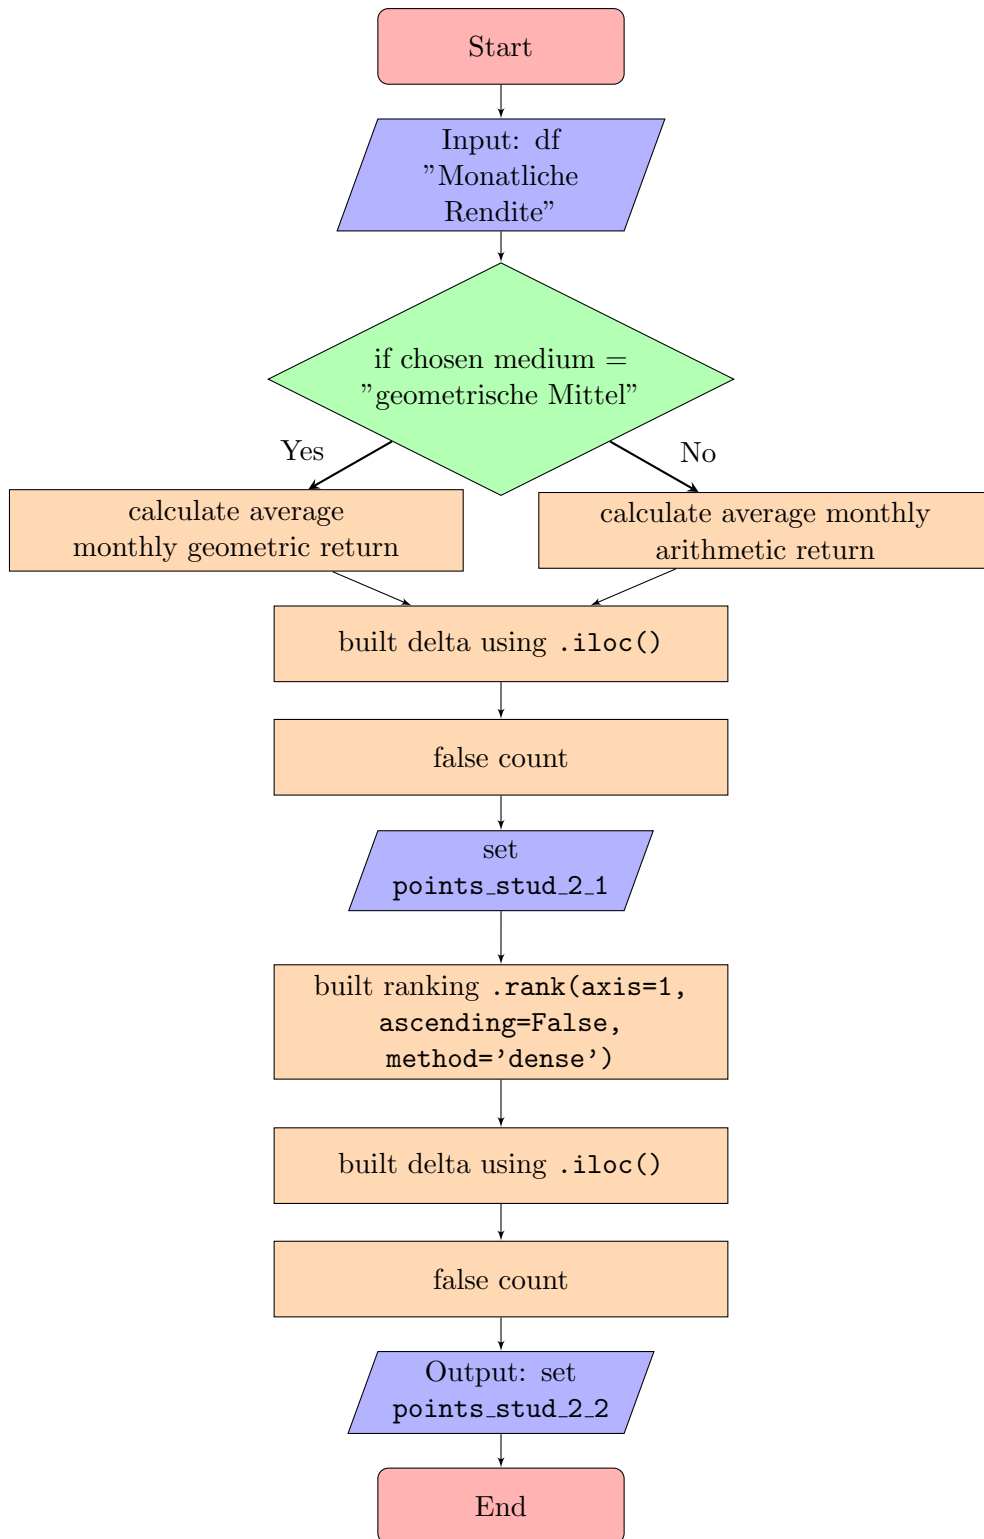
Figure 8: Flowchart Sheet Ranking

As mentioned before, the students could choose if they rather wanted to base the ranking on the arithmetic or geometric mean. The `if` statement, shown in Figure 8 above, checks whether the chosen medium is equal to the string "geometrische Mittel".

If this is deemed true, the average monthly return of the 18 stocks based on the student's individual look-back period is calculated. For that the `.rolling()` function is applied which uses the plus one return data (see Figure 7). This method provides rolling windows over the data. The size of the window is passed as a parameter in the function `.rolling(window)` and is set to the look-back period. The `.apply()` method is used in conjunction with the `.rolling()` function to apply the `np.prod` function to each rolling window of the dataframe, which calculates the product of all the values within the window. The `raw=True` parameter indicates that the function is applied to the underlying numpy array of the dataframe rather than to each column independently. The result is then passed to the `nthRoot` function which calculates the "n^{th}" root of the product where `n` is equal to look-back period. Finally, -1 is subtracted from the result since the returns were initially added by one. This creates a new dataframe that represents the geometric mean.

If `mittel_ranking` is not equal to "geometrische Mittel", the code inside the else block is executed. This is the case when the student has chosen the arithmetic mean. The `mean` function is then used to calculate the mean of all values in the window for each row. This creates a new dataframe that represents the arithmetic mean of `df_berechnung_mon_renditen_stud` over a rolling window of the size `lookback_period`.

Note that all calculations, except for the monthly return, are based on the student's values in order to include and reward consequential errors also with points.

After generating the solution, the delta is calculated by rounding both the student and solution dataframes to four decimal points. However, the rolling function can cause empty rows with irrelevant information, such as "look-back period". The `.iloc` function is used to delete these rows from both dataframes before calculating the delta.

As a result of deleting some rows in the dataframes, the number of points that will be subtracted for each false value needs to be modified accordingly. This adjustment ensures that the score accurately reflects the student's performance in the task even though some rows were deleted during the analyzis. The number of false values in the delta is used to calculate the points for exercise 2.1.

Based on the total return, the ranking by applying the `rank` function to `df_ranking_stud_lb` is built. The `rank` function assigns a rank to each element in a dataframe based on the values of the elements. The argument `axis` is set to `1`, meaning that the ranking will be done along the columns (across the row axis). The argument `ascending` is set to `False`, meaning that the ranking will be done in descending order. The argument method is set to `dense`, meaning that the ranks will be assigned with no gaps, i.e. if two elements have the same value they will receive the same rank and the next rank will be incremented by 2. The resulting dataframe `df_ranking_sol_rank` will have the same number of rows as `df_ranking_stud_lb` and the same number of columns but each value in the dataframe will be replaced with its rank. By building the delta, the same amount of rows are deleted simultaneously, and the false values which serve as the basis for determining the points for exercise 2.2 are counted.
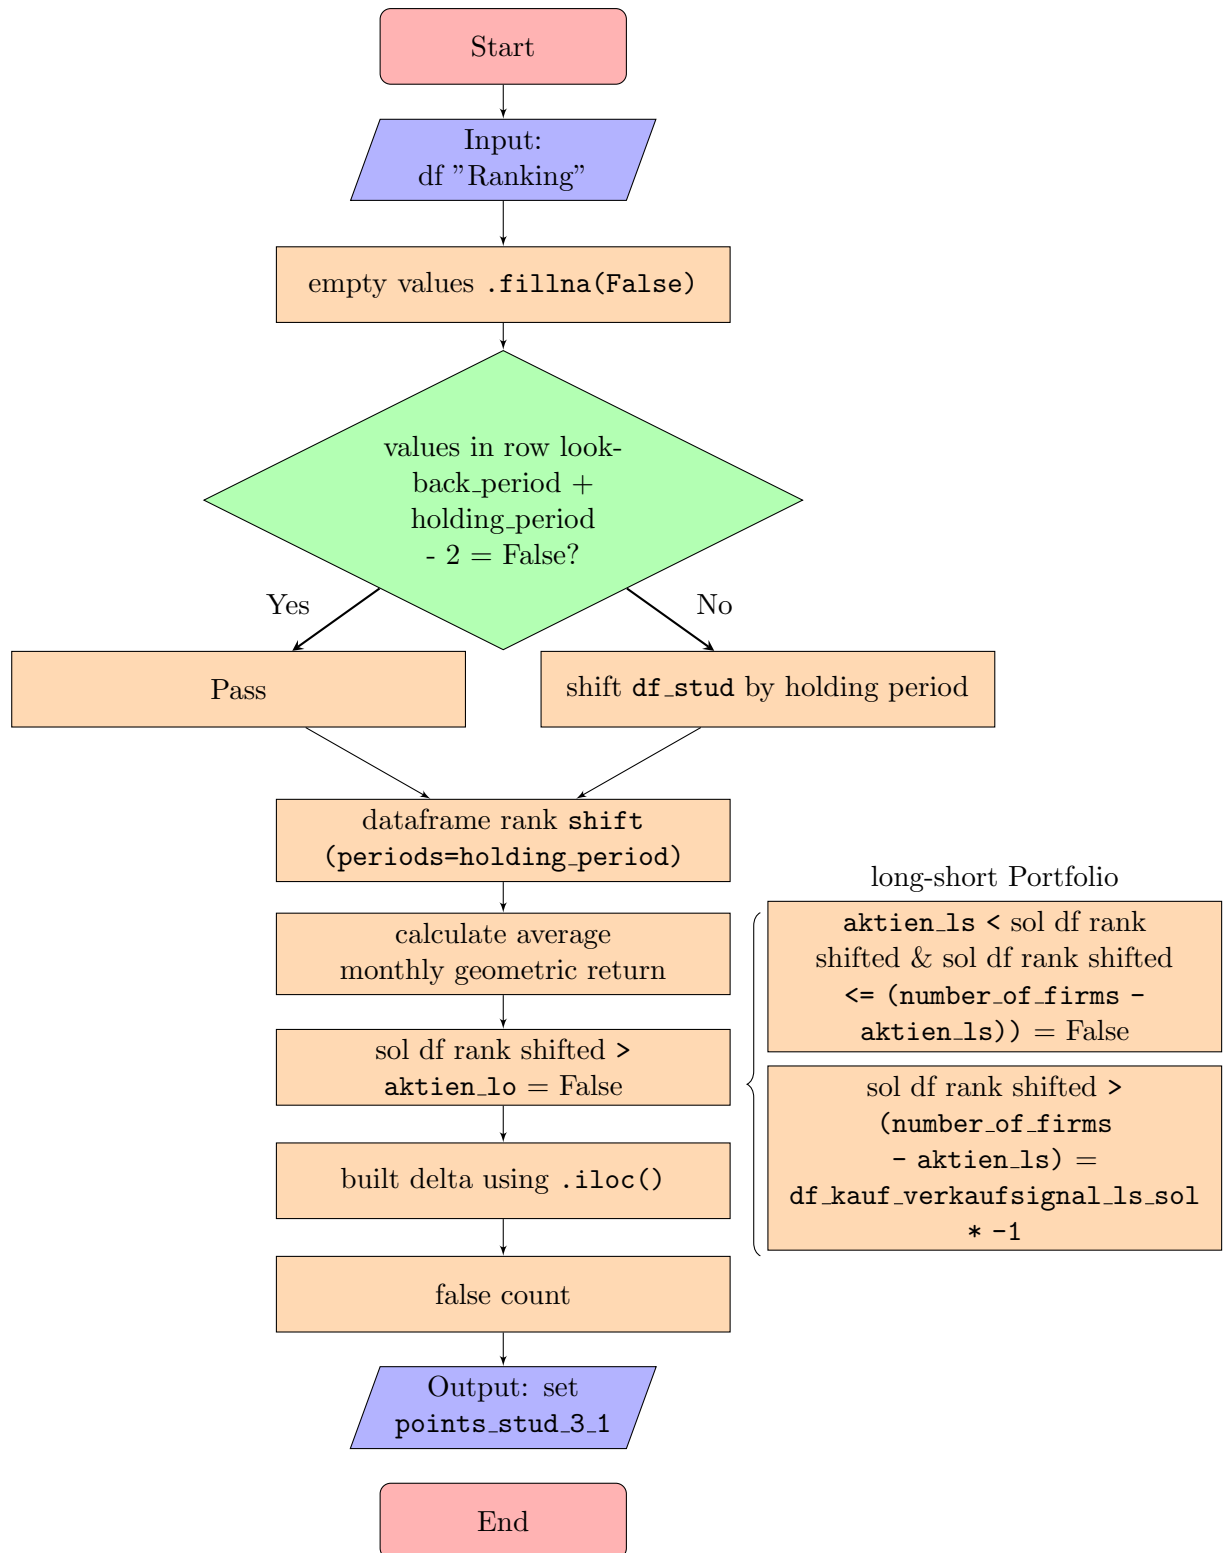
Figure 9: Flowchart Sheet Kauf- & Verkaufsignal

In Figure 9, displayed above, the students invested in the number of stocks according to their long-only or long-short strategy and calculate the total return over their holding period. Some students performed a backward test, meaning they calculated the total

return at the beginning of the holding period instead of at the end. To determine whether a backward test was conducted or not, the `.fillna(False)` function was used to replace any missing values with `False`. Then the code checks whether all the values in the row `lookback_period + holding_period - 2` of `df_kauf_verkaufsignal_lo_stud` are `False` or not. In case there is at least one `True` value in the row `lookback_period + holding_period - 1` it is definite that a backward test was performed. If there was no `True` value, the pass statement was used as a placeholder. If a backward test was performed, the code proceeds to the else statement. In that instance, in order to compare the results later, the values in the student's dataframe were shifted (using the `shift` method of the dataframe) by the holding period.

Then the `.iloc()` function is used again to delete the rows that aren't necessary for the correction. Afterwards, the dataframe with the ranks in it is shifted by the amount of holding periods so that the correct ranks will get compared with the number of stocks in the long-only portfolio for calculating the rolling return. Again by using the `.rolling()` function the moving average monthly geometric return will be calculated with the window set to the holding period.

For the long-only portfolio, the students were instructed to invest equally weighted (1/N) in the number of stocks according to their long-only strategy. Therefore, the number of stocks the student has chosen for the long-only portfolio is compared with the rank dataframe, and when the rank was greater than the number of stocks, it was replaced with `False`. After that, empty rows were deleted from both the solution and student dataframes using the `.iloc()` function, and all values were rounded to four decimal points. The delta was calculated again to compare the solution and student dataframes, and the number of false values in the delta was used to set the points for exercise 3.1.

The process is the same for the long-short portfolio except that the students buy and sell the number of shares according to their long-short strategy. Therefore, the code filters out any firms for which the rank is between the number of stocks in the long-short portfolio and the total number of firms minus the number of stocks in the long-short portfolio. This is the range of ranks that would be included in the long-short portfolio. Any stocks outside this range are set to `False` in the `df_kauf_verkaufsignal_ls_sol` dataframe indicating they should not be included in the portfolio. Because the worst-ranked stocks are sold according to the long-short portfolio the returns thereof must be multiplied by `-1`. This computation reflects the act of selling them in the long-short portfolio.
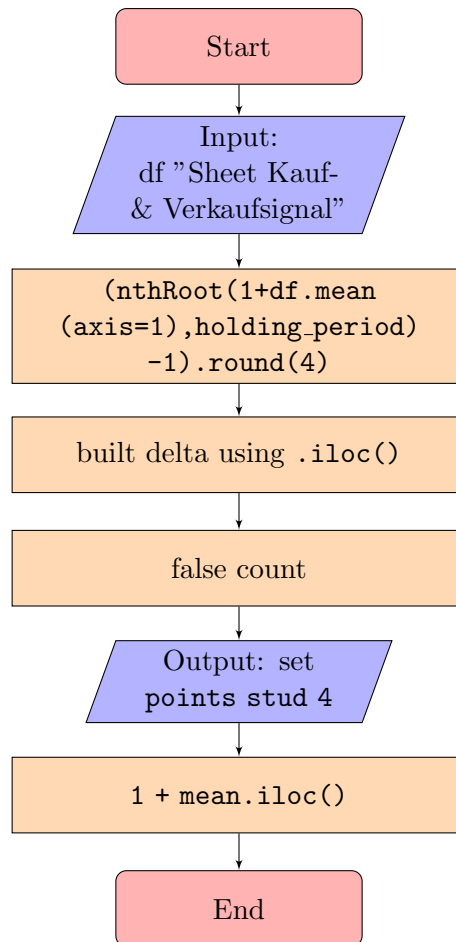
Figure 10: Flowchart Sheet Monatliche Portfoliorendite

In the presented Figure 10, the monthly portfolio returns were calculated for three different strategies, namely long-only, short-long, and buy-and-hold. The `df_kauf_verkau-fsignal_lo_stud.mean(axis=1)` calculates the mean of each row (`axis=1`) in the dataframe and returns a new dataframe with a single column. The calculation for the long-only and long-short portfolio then takes the n^{th} root of (`1 + mean returns`) with `n` set to the holding period and subtracts `1`. The final result is rounded to four decimal places using the round function. The result is stored in a new column in the `df_monatliche_portfoliorenditen_so` dataframe. For the buy-and-hold strategy, the `.mean()` function was applied to calculate the monthly returns. After rounding the values and deleting the empty rows, the deltas were computed, and the points for exercise 4 were set. The monthly return of each strategy was increased by one and the empty rows were removed again using the `.iloc()` function.
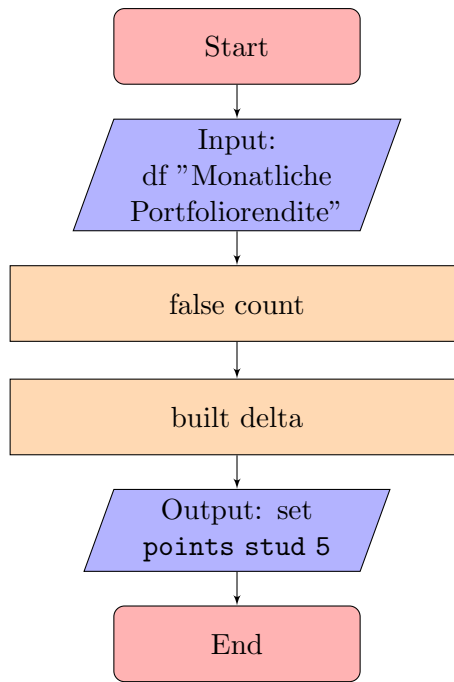
Figure 11: Flowchart Sheet Gesamtrendite und SR

In this sheet, shown in Figure 11 above, different risk and return measurements are calculated. The monthly arithmetic average return is computed using the `.mean()` function for each of the three investment strategies. Because the calculation should base on the student's value, the annualized arithmetic return is obtained using the `.iloc()` function to select a specific column from the dataframe `df_gesamtrendite_sr_stud`. To calculate the geometric return, a new dataframe is created by adding one to the monthly portfolio return of the student for each strategy. The `.iloc()` function is used to select the relevant data, and the monthly geometric average return was computed using the `nthRoot` function and the `np.prod` method. The annualized geometric return is computed using the `.iloc()` function again to select a specific column from the dataframe.

The monthly variance is calculated using the `np.var` function with the `ddof` parameter set to `1`. This parameter sets the divisor used in the calculation (`N - ddof`) where `N` represents the number of elements. By setting the `ddof` to `1`, the variance is computed in the same way as the students use the function VAR.S in Excel. The yearly variance is computed using the `.iloc()` function to select the relevant data.

The monthly and yearly volatility is obtained using the `nthRoot` function and the `.iloc()` function to select the specific row of the data. The Sharpe ratio is computed using the standard formula and plotted in the data.

All the risk and return measurements are combined into a new dataframe with an index of 1 to 9 to enable the comparison with the student's dataframe. The dataframes are rounded to four decimal points, and the delta is computed by subtracting the student's dataframe from the solution. The number of false values in the delta is counted, and

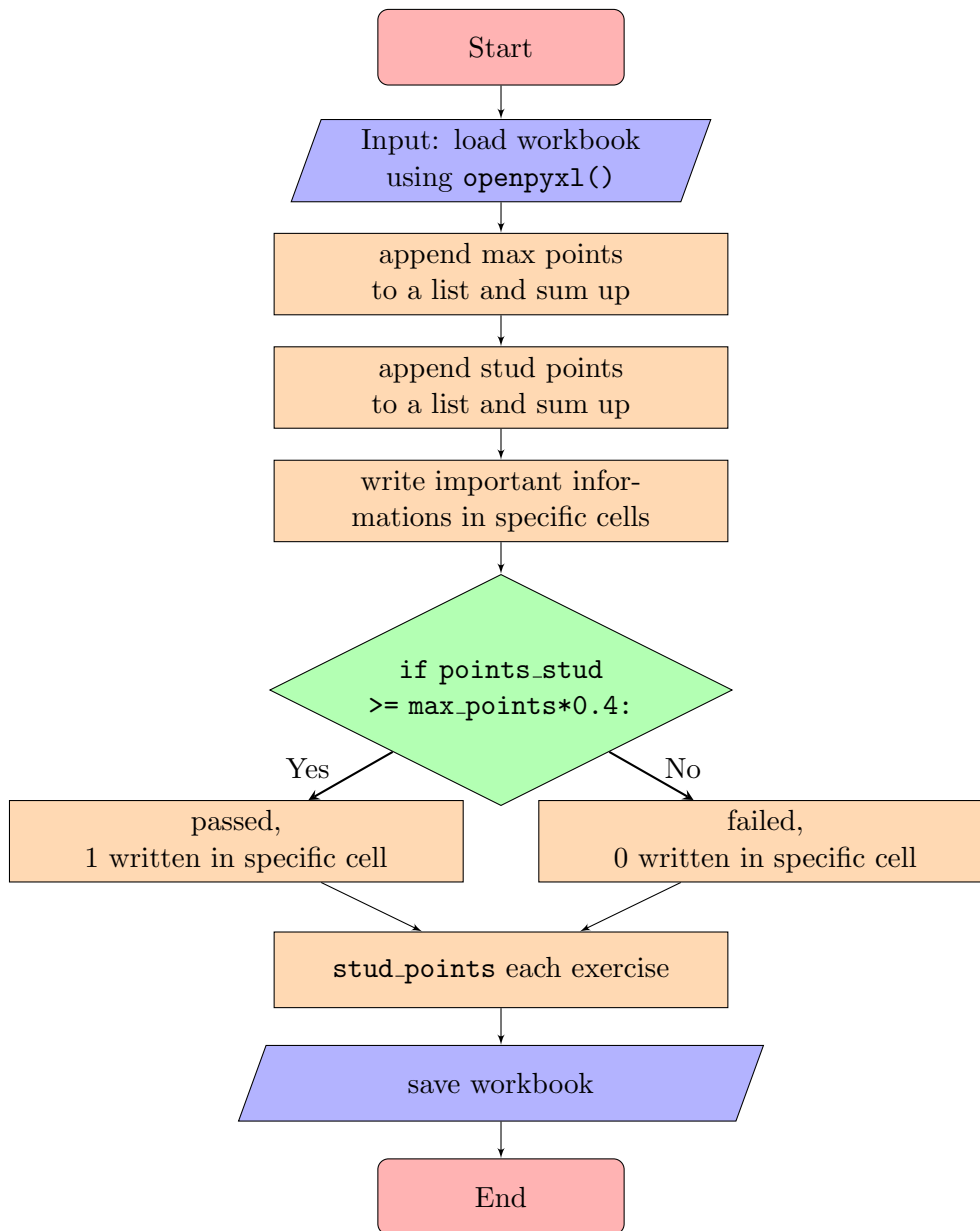the points for this exercise are set to 5.

3.3.7 Generate IA Output

Figure 12: Flowchart Sheet IA Output

After calculating and correcting the necessary dataframes, the student information and points are loaded into an empty Excel file named "IA Output". The maximum achievable points for each exercise are stored in a list, and the total number of achievable points is calculated by taking the sum of this list. The student's achieved points are calculated in the same manner. Next, the relevant student information such as first and last name, number of achieved points, and OLAT name are written into the worksheet by selecting the appropriate cell and defining the row and column in which the data should be noted.

The row of the cell is determined by additionally adding the `stud_number` variable which keeps track of the current student being processed and prevents overwriting the previous data in the Excel worksheet. To determine whether the student passed or failed the *Involving Activity 3*, an `if-else` statement is used. If the student passed the exercise, a 1 is written in the corresponding cell of the workbook. Otherwise, a 0 is noted. This simplifies the process of uploading the results to OLAT. In addition, the code records the number of points the student achieved in each individual exercise. Finally, the workbook is saved, and the correction process is repeated for the next student's file.

3.4 Correction Manual

A manual for the *Headcoach*, who is responsible for correcting the *Involving Activity 3*, is provided on Github.

4 Conclusion and further implementations

The main objective of this thesis was to develop an automated correction tool to efficiently correct the *Involving Activity 3* in the course *Asset Management: Investments* at the University of Zurich. To ensure flexibility with regard to look-back and holding periods, an option for individual selection has been incorporated into the *Involving Activity 3* at the outset of the exercise. Additionally, the input data comprising historical stock prices of 18 shares has been loaded into Python. As a result, the correction tool designed in Python is capable of accommodating updates to input data, such as changes in share prices or dates, and also on the look-back and holding period. Moreover, in order to provide the students with more flexibility in determining their rankings, an option was implemented where students are allowed to choose between using either the arithmetic or geometric mean as the basis for their ranking. This would allow students to select the method that best aligns with their personal preferences and understanding of the data. The correction tool is designed to consider consequential errors by using the student's submitted values for correction except for the first exercise for calculating the monthly returns.

The tool has successfully automated the correction of the *Involving Activity 3*, obviating the need for manual correction. The *Headcoach* is only required to make certain modifications in the code, such as adjusting the filepath (see Readme on Github). The tool provides the capability to adjust the loaded workbook range for each sheet or the saved input variables (e.g. the number of firms) in the event that modifications are required.

Nevertheless, some potential limitations must be considered. Firstly, because of loading the range of every workbook into a dataframe in Python, the range is predetermined. One potential solution to this limitation is to automate the process of adjusting the range. This could be achieved by implementing a function that can read the active range, as a number of firms and time period, in Excel based on certain criteria, such as colour-coding or specific cell values. Additionally, it may be possible to optimize the code to handle larger datasets more efficiently, thereby reducing the need for manual adjustments. However, it is important to note that these solutions may require additional development and testing to ensure their effectiveness and reliability.

Secondly, to address the drawback of not being able to provide feedback to students, a possible extension of the tool could include generating personalised feedback for each student, such as a sheet with a point overview or their Excel sheet color-coded that highlights correct and incorrect answers. This would enhance the usefulness of the tool by providing students with valuable feedback on their performance, enabling them to identify areas for improvement and promoting a more effective learning experience.

However, the use of the correction tool remains an effective and efficient approach for correcting the *Involving Activity 3*. The objective and automated nature of the correction process offers several distinct advantages including:

- significant reduction of processing times

- lower incidence of errors

- increased consistency in grading

- enhanced fairness in assessment

- facilitation of cross-comparisons among students

Concluding, the correction tool is a valuable addition to the course *Asset Management: Investments.*

Bibliography

Asness, C. S., Moskowitz, T. J., and Pedersen, L. H. (2013). Value and momentum everywhere. *The Journal of Finance*, 68(3):929–985.

Baltas, A.-N. and Kosowski, R. (2013). Momentum strategies in futures markets and trend-following funds. *Available at SSRN 1968996*.

Barberis, N., Shleifer, A., and Vishny, R. (1998). A model of investor sentiment. *Journal of Financial Economics*, 49(1):307–343.

Beyhaghi, M. and Ehsani, S. (2016). The cross-section of expected returns in the secondary corporate loan market. *The Review of Asset Pricing Studies*, 7(2):243–277.

Bhattacharya, D., Li, W.-H., and Sonaer, G. (2015). Has momentum lost its momentum? *Review of Quantitative Finance and Accounting*, 48(1):191–218.

Bhojraj, S. and Swaminathan, B. (2006). Macromomentum: Returns predictability in international equity indices. *The Journal of Business*, 79(1):429–451.

Carhart, M. M. (1997). On persistence in mutual fund performance. *The Journal of Finance*, 52(1):57–82.

Chan, K., Hameed, A., and Tong, W. (2000). Profitability of momentum strategies in the international equity markets. *The Journal of Financial and Quantitative Analysis*, 35(2):153–172.

Chan, L. K., Jegadeesh, N., and Lakonishok, J. (2019). The profitability of momentum strategies. *Financial Analysts Journal*, 55(6):80–90.

Cheemaa, M. A., Narteab, G. V., and Szulczyk, K. R. (2018). Cross-sectional and time-series momentum returns and market dynamics: evidence from japan. *Applied Economics*, 50(23):2600–2612.

Chui, A. C., Titman, S., and Wei, K. J. (2010). Individualism and momentum around the world. *The Review of Financial Studies*, 65(1):361–392.

D'Souza, I., Srichanachaichok, V., Wang, G., and Yao, Y. (2016). The enduring effect of time-series momentum on stock returns over nearly 100-years. *Available at SSRN 2720600*.

Ehsani, S. and Linnainmaa, J. A. (2022). Factor momentum and the momentum factor. *The Journal of Finance*, 77(3):1877–1919.

Fama, E. F. (1970). Efficient capital markets: A review of theory and empirical work. *Journal of Finance*, 25(2):383–417.

Fang, J., Hao, W., and Wongchoti, U. (2021). Time-series momentum in individual stocks: is it there and where to look? *Applied Economics*, 54(18):2048–2066.

Gao, L., Han, Y., Li, S. Z., and Zhou, G. (2018). Market intraday momentum. *Journal of Financial Economics*, 129(2):394–414.

Georgopoulou, A. and Wang, J. G. (2017). The trend is your friend: Time- series momentum strategies across equity and commodity markets. *Review of Finance*, 21(4):1557–1592.

Goyal, A. and Jegadeesh, N. (2017). Cross-sectional and time-series tests of return predictability: What is the difference? *The Review of Financial Studies*, 31(5):1784–1824.

Griffin, J. M., Ji, X., and Martin, J. S. (2005). Global momentum strategies. *The Journal of Portfolio Management*, 31(2):23–39.

Grinblatt, M. and Han, B. (2002). The disposition effect and momentum. Working Paper 8734, National Bureau of Economic Research.

Grinblatt, M., Titman, S., and Wermers, R. (1995). Momentum investment strategies, portfolio performance, and herding: A study of mutual fund behavior. *The American Economic Review*, 85(5):1088–1105.

Ham, H., Cho, H., Kim, H., and Ryu, D. (2019). Timeâseries momentum in china's commodity futures market. *Journal of Futures Markets*, 39(12):1515–1528.

Hameed, A. and Kusnadi, Y. (2002). Momentum strategies: Evidence from pacific basin stock markets. *The Journal of Financial Research*, 25(3):383–397.

He, X.-Z. and Li, K. (2015). Profitability of time series momentum. *Journal of Banking and Finance*, 53:140–157.

Hoitash, R. and Krishnan, M. M. (2007). Herding, momentum and investor over-reaction. *Review of Quantitative Finance and Accounting*, 30:25–47.

Hong, H. and Stein, J. C. (1999). A unified theory of underreaction, momentum trading, and overreaction in asset markets. *The Journal of Finance*, 54(6):2143–2184.

Huang, D., Li, J., Wang, L., and Zhou, G. (2020). Time series momentum: Is it there? *Journal of Financial Economics*, 135(3):774–794.

Hurst, B., Ooi, Y. H., and Pedersen, L. H. (2017). A century of evidence on trend-following investing. *The Journal of Portfolio Management*, 44(1):15–29.

Jegadeesh, N. and Titman, S. (1993). Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of Finance*, 48(1):65–91.

Jin, M., Kearney, F., Li, Y., and Yang, Y. C. (2019). Intraday timeâseries momentum: Evidence from china. *Journal of Investment Management*, 40(4):632–650.

Jostova, G., Nikolova, S., Philipov, A., and Stahel, C. W. (2019). Momentum in corporate bond returns. *The Review of Financial Studies*, 26(7):1649–1693.

Kang, J., Liu, M.-H., and Xiaoyan Ni, S. (2002). Contrarian and momentum strategies in the china stock market: 1993â2000. *Pacific-Basin Finance Journal*, 10(3):243–265.

Kim, A. Y., Tse, Y., and Wald, J. K. (2016). Time series momentum and volatility scaling. *Journal of Financial Markets*, 30:103–124.

Levy, R. A. (1967). Relative strength as a criterion for investment selection. *The Journal of Finance*, 22(4):595–610.

Lewellen, J. (2002). Momentum and autocorrelation in stock returns. *The Review of Financial Studies*, 15(2):533–564.

Lim, B. Y., Wang, J. G., and Yao, Y. (2018). Time-series momentum in nearly 100 years of stock returns. *Journal of Banking and Finance*, 97:283–296.

Menkhoff, L., Sarno, L., Schmeling, M., and Schrimpf, A. (2012). Currency momentum strategies. *Journal of Financial Economics*, 106(3):660–684.

Miffre, J. and Rallis, G. (2007). Momentum strategies in commodity futures markets. *Journal of Banking and Finance*, 31(6):1863–1886.

Moskowitz, T. J. and Grinblatt, M. (1999). Do industries explain momentum? *The Review of Financial Studies*, 54(4):1249–1290.

Moskowitz, T. J., Ooi, Y. H., and Pedersen, L. H. (2012). Time series momentum. *Journal of Finance*, 104(2):228–250.

Naughton, T., Truong, C., and Veeraraghavan, M. (2008). Momentum strategies and stock returns: Chinese evidence. *Pacific-Basin Finance Journal*, 16(4):476–492.

Okunev, J. and White, D. (2003). Do momentum-based strategies still work in foreign currency markets? *Journal of Financial and Quantitative Analysis*, 38(2):425–447.

Rouwenhorst, K. G. (1998). International momentum strategies. *The Journal of Finance*, 53(1):267–284.

Shi, H.-L. and Zhou, W.-X. (2017). Time series momentum: Is it there? *Physica A*, 483:309–318.

Zakamulin, V. and Giner, J. (2019). Trend following with momentum versus moving averages: a tale of differences. *Quantitative Finance*, 20(6):985–1007.

Appendices

A   Source Code

The source code developed and used during this thesis can be found on: `https: //github.com/saskse/Testing_Momentum_Strategies_using_Python_BT`

B   Statutory Declaration

I hereby declare that the thesis with title

"Testing Momentum Strategies using Python"

has been composed by myself autonomously and that no means other than those declared were used. In every single case, I have indicated parts that were taken out of published or unpublished work, either verbatim or in a paraphrased manner, as such through a quotation.

This thesis has not been handed in or published before in the same or similar form.

Zurich, March 9, 2023

———————————
Saskia Senn